

## part3

October 20, 2025

```
[8]: # =====
# ABC Hotels - Booking Cancellation Prediction (Final, Condensed)
# Reproducible notebook: loads data, engineers features,
# compares 2 dense NNs, plots ROC & calibration, learning curve,
# computes permutation importance, saves artifacts.
# =====
import warnings, os
warnings.filterwarnings("ignore")

from pathlib import Path
import numpy as np, pandas as pd, matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split, learning_curve
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import OneHotEncoder, StandardScaler
from sklearn.impute import SimpleImputer
from sklearn.neural_network import MLPClassifier
from sklearn.metrics import (
    roc_auc_score, accuracy_score, precision_recall_fscore_support,
    confusion_matrix, classification_report, roc_curve
)
from sklearn.calibration import calibration_curve
from sklearn.inspection import permutation_importance

# ----- Config -----
CSV_PATH = "project_data.csv"
RANDOM_STATE = 21
np.random.seed(RANDOM_STATE)

ARTIFACTS_DIR = Path("artifacts"); ARTIFACTS_DIR.mkdir(parents=True, exist_ok=True)
FIG_DIR = ARTIFACTS_DIR / "figures"; FIG_DIR.mkdir(parents=True, exist_ok=True)

# ----- Load & Prepare -----
df = pd.read_csv(CSV_PATH)
print(f"Loaded dataset with {len(df)} rows and {len(df.columns)} columns.")
```

```

if "Booking_ID" in df.columns:
    df = df.drop(columns=["Booking_ID"])

df["booking_status"] = df["booking_status"].astype(str).str.strip().str.lower().
    map({
        "canceled":1, "cancelled":1,
        "not_canceled":0, "not cancelled":0, "not-canceled":0, "not canceled":0
    })
df = df.dropna(subset=["booking_status"])
df["booking_status"] = df["booking_status"].astype(int)

if "arrival_date" in df.columns:
    df["arrival_date"] = pd.to_datetime(df["arrival_date"], errors="coerce")
    df["arrival_month"] = df["arrival_date"].dt.month
    def _m2s(m):
        if pd.isna(m): return np.nan
        m = int(m)
        return "winter" if m in [12,1,2] else "spring" if m in [3,4,5] else
            "summer" if m in [6,7,8] else "fall"
    df["arrival_season"] = df["arrival_month"].map(_m2s)

y = df["booking_status"]
X = df.drop(columns=["booking_status"])

cat_cols = [c for c in X.columns if X[c].dtype == "object"]
for c in [
    "type_of_meal_plan", "room_type_reserved", "market_segment_type", "arrival_season"]:
    if c in X.columns and c not in cat_cols and not pd.api.types.
        is_numeric_dtype(X[c]):
            cat_cols.append(c)
num_cols = [c for c in X.columns if c not in cat_cols and pd.api.types.
    is_numeric_dtype(X[c])]

numeric_tf = Pipeline([
    ("imputer", SimpleImputer(strategy="median")),
    ("scaler", StandardScaler())
])
onehot_args = {"handle_unknown": "ignore"}
if "sparse_output" in OneHotEncoder.__init__.code.co_varnames:
    onehot_args["sparse_output"] = False
else:
    onehot_args["sparse"] = False
categorical_tf = Pipeline([
    ("imputer", SimpleImputer(strategy="most_frequent")),
    ("onehot", OneHotEncoder(**onehot_args))
])

```

```

preprocessor = ColumnTransformer([('num', numeric_tf, num_cols), ("cat", categorical_tf, cat_cols)], remainder="drop")

X_train, X_valid, y_train, y_valid = train_test_split(X, y, test_size=0.20, stratify=y, random_state=RANDOM_STATE)

# ----- Models -----
mlp_A = MLPClassifier(hidden_layer_sizes=(64,32), activation="relu",
    solver="adam",
    alpha=1e-4, batch_size=512, learning_rate="adaptive",
    learning_rate_init=1e-3, max_iter=50, early_stopping=True,
    validation_fraction=0.15, n_iter_no_change=5,
    random_state=RANDOM_STATE, verbose=False)

mlp_B = MLPClassifier(hidden_layer_sizes=(32,16), activation="relu",
    solver="adam",
    alpha=5e-4, batch_size=256, learning_rate="adaptive",
    learning_rate_init=1e-3, max_iter=75, early_stopping=True,
    validation_fraction=0.15, n_iter_no_change=5,
    random_state=RANDOM_STATE, verbose=False)

pipe_A = Pipeline([('preprocessor', preprocessor), ("clf", mlp_A)])
pipe_B = Pipeline([('preprocessor', preprocessor), ("clf", mlp_B)])

print("\nTraining Model A (64,32)..."); pipe_A.fit(X_train, y_train);
print("Model A trained.")
print("\nTraining Model B (32,16)..."); pipe_B.fit(X_train, y_train);
print("Model B trained.")

# ----- Evaluation -----
def evaluate_model(pipeline, X_tr, y_tr, X_va, y_va, label="Model"):
    y_tr_proba = pipeline.predict_proba(X_tr)[:,1]
    y_va_proba = pipeline.predict_proba(X_va)[:,1]
    y_tr_pred = (y_tr_proba>=0.5).astype(int); y_va_pred = (y_va_proba>=0.5).astype(int)
    M = {}
    M["train_auc"] = roc_auc_score(y_tr, y_tr_proba)
    M["valid_auc"] = roc_auc_score(y_va, y_va_proba)
    M["train_acc"] = accuracy_score(y_tr, y_tr_pred)
    M["valid_acc"] = accuracy_score(y_va, y_va_pred)
    p,r,f1,_ = precision_recall_fscore_support(y_va, y_va_pred,
        average="binary", zero_division=0)
    M["valid_precision"], M["valid_recall"], M["valid_f1"] = p,r,f1
    M["confusion_matrix"] = confusion_matrix(y_va, y_va_pred)
    M["classification_report"] = classification_report(y_va, y_va_pred,
        target_names=["Not Canceled", "Canceled"])

```

```

print(f"\n==== {label} ====")
print(f"Train AUC: {M['train_auc']:.3f} | Valid AUC: {M['valid_auc']:.3f}")
print(f"Train Acc: {M['train_acc']:.3f} | Valid Acc: {M['valid_acc']:.3f}")
print(f"Precision: {M['valid_precision']:.3f} | Recall: {M['valid_recall']:.3f} | F1: {M['valid_f1']:.3f}")
print("Confusion Matrix:\n", M["confusion_matrix"])
print("Classification Report:\n", M["classification_report"])
return M, y_va_proba

metrics_A, y_valid_proba_A = evaluate_model(pipe_A, X_train, y_train, X_valid,
                                             y_valid, "Model A (64,32)")
metrics_B, y_valid_proba_B = evaluate_model(pipe_B, X_train, y_train, X_valid,
                                             y_valid, "Model B (32,16)")

pd.DataFrame([
    {"model": "A (64,32)", **{k:v for k,v in metrics_A.items() if k not in ["confusion_matrix", "classification_report"]}},
    {"model": "B (32,16)", **{k:v for k,v in metrics_B.items() if k not in ["confusion_matrix", "classification_report"]}},
]).to_csv(ARTIFACTS_DIR/"metrics_summary.csv", index=False)

# ----- ROC (A vs B) -----
fpr_A, tpr_A, _ = roc_curve(y_valid, y_valid_proba_A)
fpr_B, tpr_B, _ = roc_curve(y_valid, y_valid_proba_B)
plt.figure(figsize=(6,5))
plt.plot(fpr_A, tpr_A, label=f"Model A (AUC={metrics_A['valid_auc']:.3f})")
plt.plot(fpr_B, tpr_B, label=f"Model B (AUC={metrics_B['valid_auc']:.3f})")
plt.plot([0,1],[0,1], '--', lw=1)
plt.xlabel("False Positive Rate"); plt.ylabel("True Positive Rate")
plt.title("ROC Curve - Validation"); plt.legend(); plt.grid(True, alpha=0.3)
plt.tight_layout(); plt.savefig(FIG_DIR/"roc_curve_A_vs_B.png", dpi=150); plt.show()

# ----- Calibration (A vs B) -----
prob_true_A, prob_pred_A = calibration_curve(y_valid, y_valid_proba_A,
                                              n_bins=10, strategy="uniform")
prob_true_B, prob_pred_B = calibration_curve(y_valid, y_valid_proba_B,
                                              n_bins=10, strategy="uniform")
plt.figure(figsize=(6,5))
plt.plot([0,1],[0,1], "--", label="Perfectly calibrated")
plt.plot(prob_pred_A, prob_true_A, marker="o", label="Model A")
plt.plot(prob_pred_B, prob_true_B, marker="o", label="Model B")
plt.xlabel("Mean predicted probability"); plt.ylabel("Fraction of positives")
plt.title("Calibration Curve - Validation"); plt.legend(); plt.grid(True, alpha=0.3)

```

```

plt.tight_layout(); plt.savefig(FIG_DIR/"calibration_curve_A_vs_B.png", dpi=150); plt.show()

# ----- Select Final & Learning Curve -----
final_model = pipe_A
print("\nSelected Final Model: Model A (64,32)")

train_sizes, train_scores, valid_scores = learning_curve(
    final_model, X, y, cv=3, scoring="roc_auc",
    train_sizes=np.linspace(0.1,1.0,5), n_jobs=-1
)
plt.figure(figsize=(6,5))
plt.plot(train_sizes, train_scores.mean(axis=1), marker='o', label="Training")
plt.plot(train_sizes, valid_scores.mean(axis=1), marker='o', label="Validation")
plt.xlabel("Training examples"); plt.ylabel("ROC-AUC"); plt.title("Learning Curve - Final Model")
plt.grid(True, alpha=0.3); plt.legend(); plt.tight_layout()
plt.savefig(FIG_DIR/"learning_curve_final_model.png", dpi=150); plt.show()

# ----- Confusion Matrix (thr=0.5) -----
y_valid_pred_final = (y_valid_proba_A >= 0.5).astype(int)
cm = confusion_matrix(y_valid, y_valid_pred_final)
plt.figure(figsize=(4.5,4.5))
plt.imshow(cm, cmap="Blues")
for i in range(2):
    for j in range(2):
        plt.text(j, i, cm[i,j], ha="center", va="center")
plt.xticks([0,1], ["Not Canceled", "Canceled"]); plt.yticks([0,1], ["Not Canceled", "Canceled"])
plt.xlabel("Predicted"); plt.ylabel("Actual"); plt.title("Confusion Matrix - Final Model (0.5)")
plt.tight_layout(); plt.savefig(FIG_DIR/"confusion_matrix_final.png", dpi=150); plt.show()

# ----- Permutation Importance -----
def get_feature_names_from_ct(ct: ColumnTransformer) -> list:
    names = []
    for name, transformer, cols in ct.transformers_:
        if name == "remainder" and transformer == "drop": continue
        if hasattr(transformer, "get_feature_names_out"):
            try: fn = transformer.get_feature_names_out(cols)
            except: fn = transformer.get_feature_names_out()
            names.extend(fn)
        elif hasattr(transformer, "steps"):
            last = transformer.steps[-1][1]
            if hasattr(last, "get_feature_names_out"):
                try: fn = last.get_feature_names_out(cols)

```

```

        except: fn = last.get_feature_names_out()
        names.extend(fn)
    else:
        names.extend(cols if isinstance(cols, list) else [cols])
    else:
        names.extend(cols if isinstance(cols, list) else [cols])
return names

final_model.fit(X_train, y_train)
perm = permutation_importance(final_model, X_valid, y_valid, scoring="roc_auc",
                               n_repeats=5, random_state=RANDOM_STATE, n_jobs=-1)
feat_names = get_feature_names_from_ct(final_model.named_steps["preprocessor"])
if len(feat_names) != len(perm.importances_mean):
    feat_names = [f"feature_{i}" for i in range(len(perm.importances_mean))]
imp_df = pd.DataFrame({"feature":feat_names,
                       "importance_mean":perm.importances_mean,
                       "importance_std":perm.importances_std}).
    ↪sort_values("importance_mean", ascending=False)
imp_df.to_csv(ARTIFACTS_DIR/"permutation_importance_validation.csv", ↪
    ↪index=False)

print("\nTop 20 features by permutation importance (AUC decrease):")
print(imp_df.head(20).to_string(index=False))

plt.figure(figsize=(8,6))
top = imp_df.head(20)[::-1]
plt.barh(top[["feature"]], top[["importance_mean"]])
plt.xlabel("Mean importance (AUC decrease)")
plt.title("Permutation Feature Importance - Final Model")
plt.tight_layout(); plt.savefig(FIG_DIR/"permutation_importance_top20.png", ↪
    ↪dpi=150); plt.show()

# ----- Save reports -----
with open(ARTIFACTS_DIR/"classification_report_model_A.txt","w") as f: f.
    ↪write(metrics_A["classification_report"])
with open(ARTIFACTS_DIR/"classification_report_model_B.txt","w") as f: f.
    ↪write(metrics_B["classification_report"])

print(f"\nAll artifacts saved to: {ARTIFACTS_DIR.resolve()}")
print("Figures saved:", [p.name for p in FIG_DIR.iterdir()])

```

Loaded dataset with 36,238 rows and 17 columns.

Training Model A (64,32)...  
Model A trained.

Training Model B (32,16)...

Model B trained.

==== Model A (64,32) ====  
Train AUC: 0.933 | Valid AUC: 0.921  
Train Acc: 0.865 | Valid Acc: 0.861  
Precision: 0.836 | Recall: 0.716 | F1: 0.771

Confusion Matrix:

```
[[4539 333]
 [ 675 1701]]
```

Classification Report:

	precision	recall	f1-score	support
Not Canceled	0.87	0.93	0.90	4872
Canceled	0.84	0.72	0.77	2376
accuracy			0.86	7248
macro avg	0.85	0.82	0.84	7248
weighted avg	0.86	0.86	0.86	7248

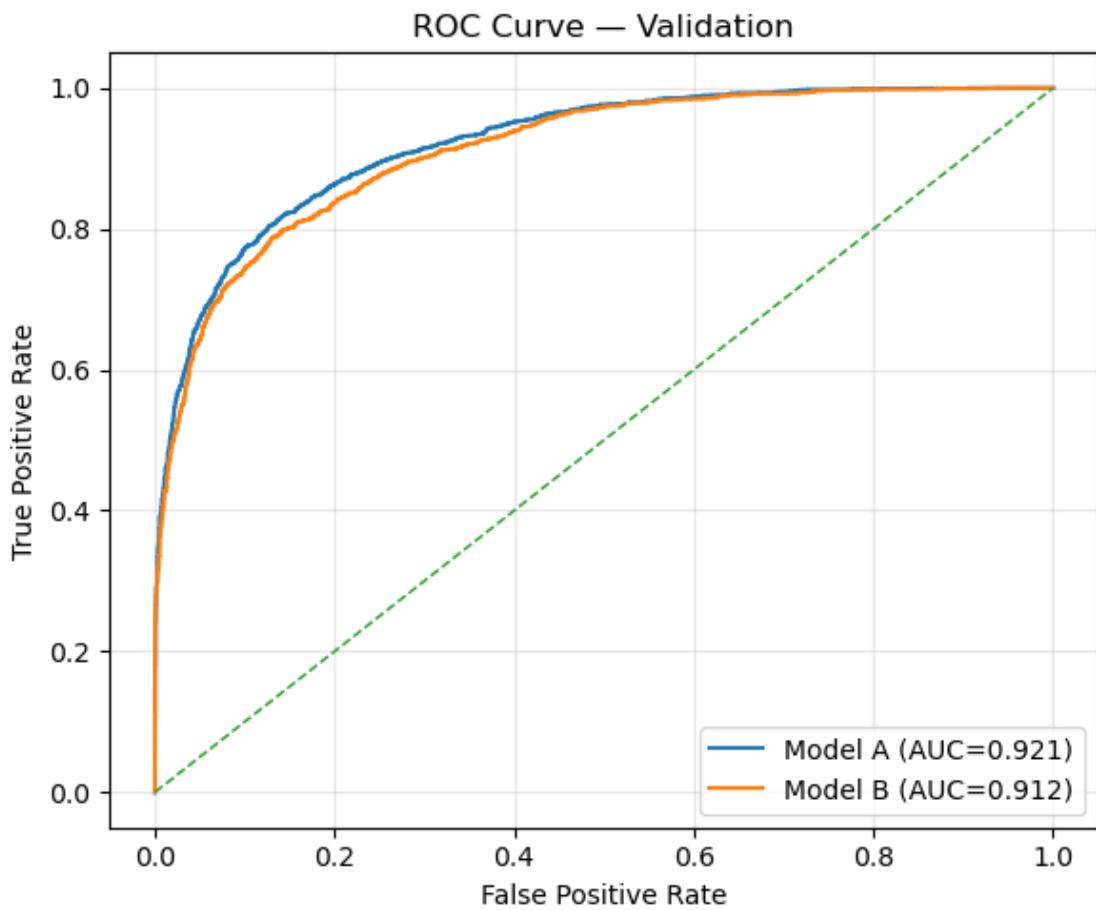
==== Model B (32,16) ====  
Train AUC: 0.915 | Valid AUC: 0.912  
Train Acc: 0.854 | Valid Acc: 0.854  
Precision: 0.812 | Recall: 0.722 | F1: 0.764

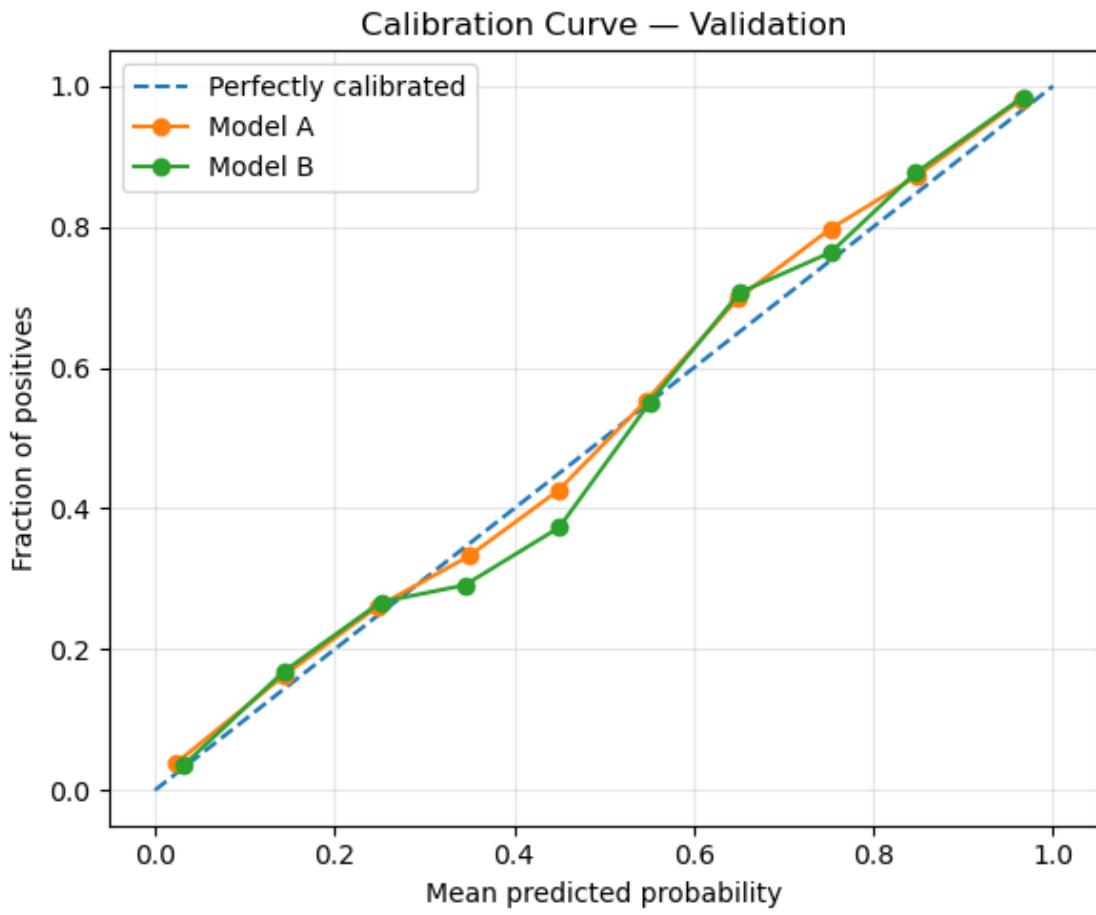
Confusion Matrix:

```
[[4474 398]
 [ 660 1716]]
```

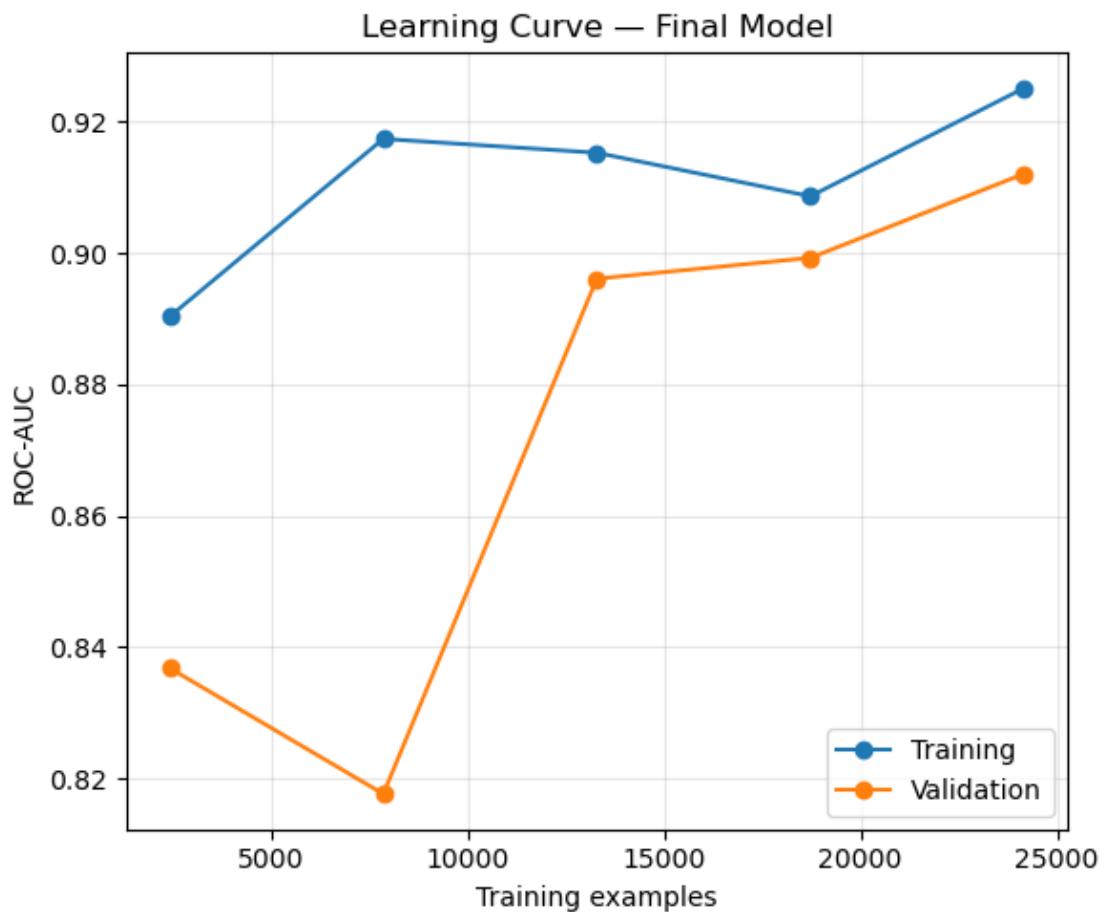
Classification Report:

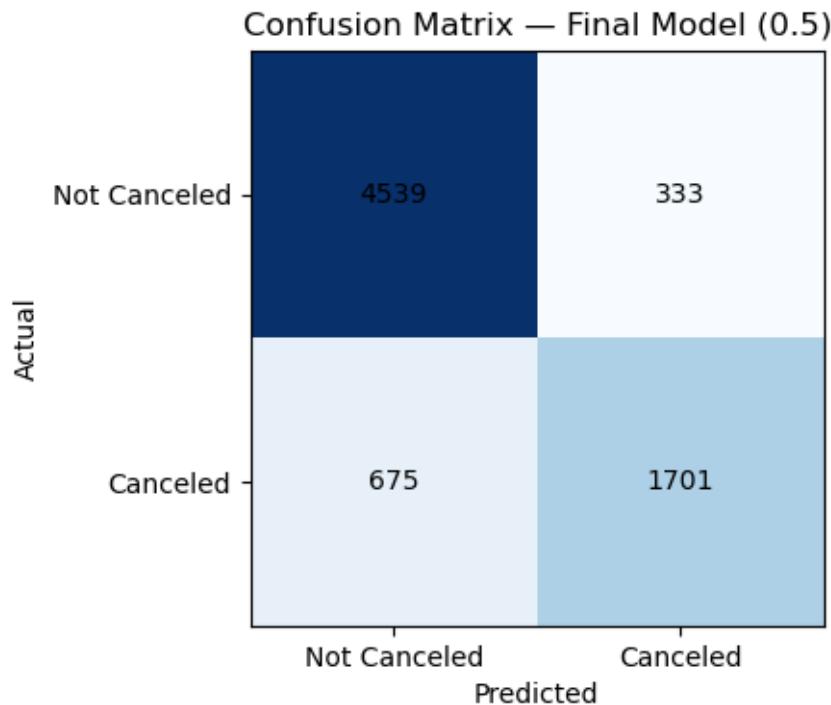
	precision	recall	f1-score	support
Not Canceled	0.87	0.92	0.89	4872
Canceled	0.81	0.72	0.76	2376
accuracy			0.85	7248
macro avg	0.84	0.82	0.83	7248
weighted avg	0.85	0.85	0.85	7248





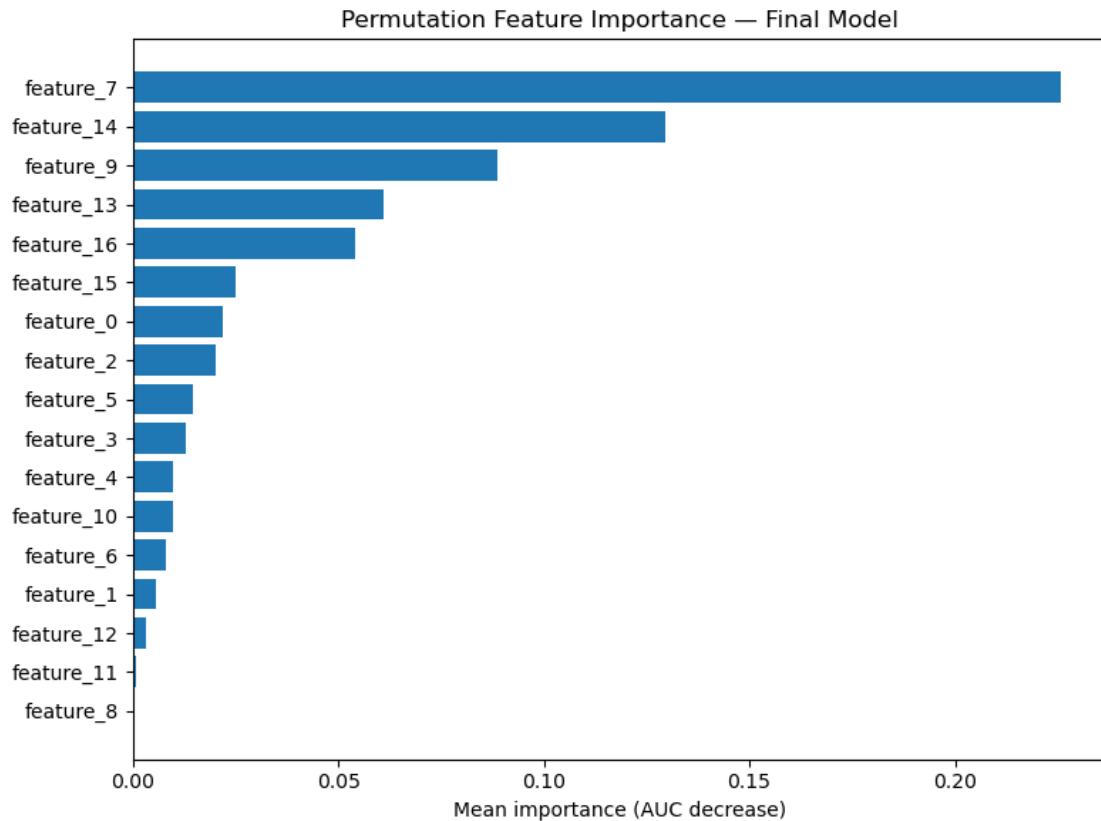
Selected Final Model: Model A (64,32)





Top 20 features by permutation importance (AUC decrease):

feature	importance_mean	importance_std
feature_7	0.225909	0.002639
feature_14	0.129432	0.005684
feature_9	0.088639	0.004287
feature_13	0.060826	0.003301
feature_16	0.053924	0.002040
feature_15	0.024885	0.000681
feature_0	0.021880	0.000766
feature_2	0.020049	0.000723
feature_5	0.014594	0.000786
feature_3	0.012806	0.001816
feature_4	0.009673	0.000349
feature_10	0.009655	0.001910
feature_6	0.007997	0.002147
feature_1	0.005421	0.000989
feature_12	0.003179	0.000925
feature_11	0.000845	0.000313
feature_8	0.000000	0.000000



```
All artifacts saved to: /Users/thehootch/Desktop/Data  
Science/merrimack/ml_ai/project/artifacts  
Figures saved: ['learning_curve_final_model.png', 'confusion_matrix.png',  
'learning_curve_auc.png', 'roc_curve_A_vs_B.png', 'confusion_matrix_final.png',  
'calibration_curve_A_vs_B.png', 'permutation_importance_top20.png',  
'roc_curve_validation.png']
```

```
[ ]:
```